


Computational Reproducibility in Archaeological Research: Basic Principles and a Case Study of Their Implementation

Ben Marwick

**Journal of Archaeological Method
and Theory**

ISSN 1072-5369

J Archaeol Method Theory
DOI 10.1007/s10816-015-9272-9



Your article is protected by copyright and all rights are held exclusively by Springer Science +Business Media New York. This e-offprint is for personal use only and shall not be self-archived in electronic repositories. If you wish to self-archive your article, please use the accepted manuscript version for posting on your own website. You may further deposit the accepted manuscript version in any repository, provided it is only made publicly available 12 months after official publication or later and provided acknowledgement is given to the original source of publication and a link is inserted to the published article on Springer's website. The link must be accompanied by the following text: "The final publication is available at link.springer.com".

Computational Reproducibility in Archaeological Research: Basic Principles and a Case Study of Their Implementation

Ben Marwick^{1,2} 

© Springer Science+Business Media New York 2016

Abstract The use of computers and complex software is pervasive in archaeology, yet their role in the analytical pipeline is rarely exposed for other researchers to inspect or reuse. This limits the progress of archaeology because researchers cannot easily reproduce each other's work to verify or extend it. Four general principles of reproducible research that have emerged in other fields are presented. An archaeological case study is described that shows how each principle can be implemented using freely available software. The costs and benefits of implementing reproducible research are assessed. The primary benefit, of sharing data in particular, is increased impact *via* an increased number of citations. The primary cost is the additional time required to enhance reproducibility, although the exact amount is difficult to quantify.

Keywords Reproducible research · Computer programming · Software engineering · Australian archaeology · Open science

Introduction

Archaeology, like all scientific fields, advances through rigorous tests of previously published studies. When numerous investigations are performed by different researchers and demonstrate similar results, we hold these results to be a reasonable approximation of a true account of past human behavior. This ability to reproduce the results of other researchers is a core tenet of scientific method, and when reproductions are successful, our field advances. In archaeology, we have a long tradition of empirical tests of reproducibility, for example, by returning to field sites excavated or surveyed by

✉ Ben Marwick
bmarwick@uw.edu

¹ Department of Anthropology, University of Washington, Seattle, WA, USA

² Center for Archaeological Science, University of Wollongong, Wollongong, Australia

earlier generations of archaeologists, and re-examining museum collections with new methods.

However, we, like many disciplines, have made little progress in testing the reproducibility of statistical and computational results, or even facilitating or enabling these tests (Ince *et al.* 2012; Peng 2011). The typical contemporary journal article describing the results of an archaeological study rarely contains enough information for another archaeologist to reproduce its statistical results and figures. Raw data are rarely openly and fully provided, perhaps due to the absence of data-sharing standards that acknowledge the sensitive nature of much of our data. Similarly, many of the decisions made in cleaning, tidying, analyzing, and visualizing the data are unrecorded and unreported. This is a problem because as computational results become increasingly common and complex in archaeology, and we are increasingly dependent on software to generate our results, we risk deviating from the scientific method if we are unable to reproduce the computational results of our peers (Dafoe 2014). A further problem is that when the methods are underspecified, it limits the ease with which they can be reused by the original author, and extended by others (Buckheit and Donoho 1995; Donoho *et al.* 2009; Schwab *et al.* 2000). This means that when a new methods paper in archaeology is published as a stand-alone account (*i.e.*, without any accompanying software), it is challenging and time-consuming for others to benefit from this new method. This is a substantial barrier to progress in archaeology, both in establishing the veracity of previous claims and promoting the growth of new interpretations. Furthermore, if we are to contribute to contemporary conversations outside of archaeology (as we are supposedly well-positioned to do, *cf.* Kintigh *et al.* (2014)), we need to become more efficient, interoperative, and flexible in our research. We have to be able to invite researchers from other fields into our research pipelines to collaborate in answering interesting and broad questions about past societies.

In this paper, I address these problems by demonstrating a research method that enables computational reproducibility for archaeology at the level of a familiar research product, the journal article (Fig. 1). First, I outline the general principles that motivate this approach (Table 1). These principles have been derived from software engineering and developed and refined over the last several years by researchers in computationally intensive fields such as genomics, ecology, astronomy, climatology, neuroscience, and oceanography (Stodden and Míguez 2014; Wilson *et al.* 2014). Although the data produced by some of these disciplines are often used by archaeologists, efforts toward improving reproducibility in these fields have seen little uptake among archaeologists. The principles are ordered by scope, such that the first principle is applicable to every archaeological publication that makes claims based on archaeological evidence, the second principle is applicable to most publications that contain quantitative results, and the third and fourth principles are most applicable to publications that report substantial and complex quantitative results. In the second part of the paper, I describe a case study of a recent archaeological research publication and its accompanying research compendium. In preparing this publication, I developed new methods for enabling the reproducibility of the computational results. I describe these methods and the specific tools used in this project to follow the general principles. While the specific tools used in this example will likely be replaced by others a few years from now, the general principles presented here are tool-agnostic, and can serve as a guide for archaeologists into the future.

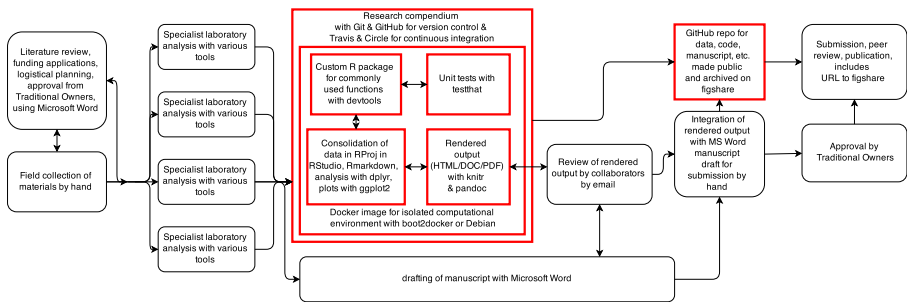


Fig. 1 Workflow diagram showing key steps and software components. The boxes with a bold outline indicate key steps and tools that enable computational reproducibility in our project

General Principles of a Reproducible Methodology

Data and Code Provenance, Sharing, and Archiving

Perhaps the most trivial principle of reproducible research is making openly available the data and methods that generated the published results. This is a computational analogue to the archaeological principle of artifact provenience. For example, without provenience information, artifacts are nearly meaningless; without providing data and code, the final published results are similarly diminished. Making data and code available enables others to inspect these materials to evaluate the reliability of the publication, and to include the materials into other projects, and may lead to higher quality and more impactful published research (Gleditsch and Strand 2003; Piwowar *et al.* 2007; Wicherts *et al.* 2011). While it might seem a basic principle for reproducible research, current community norms in archaeology, like many disciplines, do not encourage or reward the sharing of data and other materials used in the research leading to journal articles (Borgman 2012; B. McCullough 2007; Stodden *et al.* 2013; Tenopir *et al.* 2011). While funding agencies, such as the US National Science Foundation (NSF), require a data management plan (DMP) in proposals, and some journals, such as PLOS ONE and Nature, require data availability statements, none of these require all archaeologists to make their data available by default (Begley and Ioannidis 2015; Miguel *et al.* 2014). For archaeology submissions to the NSF, the DMP recommendations were developed by the Society of American Archaeologists, rather than from within the NSF (Rieth 2013).

It is difficult to prescribe a single approach to making data and other materials openly available because of the wide variety of archaeological data, and the diversity of contexts it is collected (Kintigh 2006). As a general principle that should be applicable in all cases, the provenance of the data must always be stated, even if the data are not publicly accessible (for example, due to copyright limitations, cultural sensitivities, for protection from vandalism, or because of technical limitations). Where a journal article includes data summaries and visualizations, the principle is that authors make publicly available (*i.e.*, not “by request”) the computer files containing the rawest form possible of the data from which the summaries and plots were generated (*e.g.*, spreadsheets of individual measurement records). This minimalist approach means that only the data needed to support the publication should be released, the rest can be kept private while

Table 1 Glossary of key terms used in the text

Term	Explanation	More information
Concepts		
Open source	Computer code where the source code is available for inspection, and may be freely re-used and distributed. R, Python, and GNU/Linux are all open source.	https://opensource.org/osd
Open access	Access to research products, such as publications and datasets, without financial or copyright barriers, but such that authors have control over the integrity of their work and the right to be acknowledged and cited. One approach is to publish in open access journals, such as PLOS ONE, another approach is to submit manuscripts of published papers to institutional repositories where they are freely available to the public.	http://www.budapestopenaccessinitiative.org/read
Reproducibility	A study is reproducible if there is a specific set of computational functions/analyses (usually specified in terms of code) that exactly reproduce all of the numbers and data visualizations in a published paper from raw data. Reproducibility does not require independent data collection and instead uses the methods and data collected by the original investigator.	https://osf.io/s9tya/
Replicability	A study is replicated when another researcher independently implements the same methods of data collection and analysis with a new data set.	http://languageelog.ldc.upenn.edu/nll/?p=21956
Provenance	The origin of data and code, including any transformations occurring along the way.	
File formats		
CSV	A common file format for collecting, sharing, and archiving tabular data. This is a plain text file where variables (columns) are separated by commas. Thus the name, “comma separated variables,” it is closely related to TSV, “tab separated variables.”	http://www.digitalpreservation.gov/formats/fdd/fdd000323.shtml
Plain text	A file that contains simple text characters and no formatting (<i>e.g.</i> , margins) or embedded images. Use of plain text files is not dependent on specific programs, so they can be created, read, and edited by almost any program, regardless of operating system and computer architecture. Using plain text formats allows a high degree of interoperability between computational environments, and ensures that your files can be read by other people with minimum effort. Most programming script files are plain text files.	http://www.linfo.org/plain_text.html
Binary	A file that must be interpreted by a specific program before it is human-readable and editable. For example, PDF, Microsoft Word doc, and Excel xls files are binary files, and can only be read and edited by those programs. Many commercial programs use proprietary binary file formats. This limits their interoperability and archival	

Table 1 (continued)

Term	Explanation	More information
	value. Images, video, and audio files are also binary files.	
Licenses for data and code		
CC0	Public domain, no rights reserved. This license allows for the greatest freedom for reuse. Used for data by major online repositories such as Dryad, Figshare, Zenodo. Good scientific practices assure proper credit is given <i>via</i> citation, which enforced through peer review. Marking data with CC0 sends a clear signal of zero barriers to reuse.	https://creativecommons.org/licenses/
CC-BY	Allows for reuse only if attribution is given to the author, in the manner specified by the author. Often used for copyrightable materials such as journal articles in open access publications, for example PLOS ONE, BioMed Central, and Nature Communications.	
CC-NC	Allows for reuse only for non-commercial purposes (for example, a Cultural Heritage Management business would not be allowed to use CC-NC data or code). Not recommended for most research output.	
MIT	A license especially for software that places very few restrictions on the use of the software, and disclaims the author of any responsibility for problems arising from others using the software. It is one of the most popular licenses for open source software.	http://opensource.org/licenses/MIT
Data archiving		
DOI	DOI stands for “digital object identifier,” a persistent (but not permanent) label that stores information about the online location of a electronic file. A DOI also includes metadata, for example in the case of journal article it might include the author, title, date of publication, <i>etc.</i> The online location and metadata of a file may change, but its DOI remains fixed. This means that a DOI is generally a more reliable link to an online document than a URL.	http://www.doi.org/
figshare	A commercial online digital repository where research output can be freely archived and openly accessed. Issues DOIs for individual files or groups of files.	http://figshare.com/
zenodo	Similar to figshare, but a non-profit service operated by European Organization for Nuclear Research (known as CERN).	https://zenodo.org/
tDAR	The Digital Archaeological Record (tDAR) is a digital repository for the digital records of archaeological investigations. Fees are charged for archiving files, but access to open files is free.	https://www.tdar.org/
Open Context		http://opencontext.org/

Table 1 (continued)

Term	Explanation	More information
	A data publishing and archiving service. It is aimed at maximizing the integration of data with other services (such as maps, media, and other data sets). Similar to tDAR, there are fees to upload but accessing open data is free.	
Archaeological Data Service	An open data repository focused on output from research and commercial archaeology in the UK. There are fees to upload but accessing open data is free.	http://archaeologydataservice.ac.uk/
CLOCKSS	A not-for-profit joint venture between several academic publishers and research libraries to build a sustainable, geographically distributed dark archive with which to ensure the long-term survival of Web-based scholarly publications.	https://www.clockss.org/
Document markup languages		
markdown	A simple, minimal language for formatting plain text files so that they can be converted into richly formatted HTML, PDF, and Microsoft Word documents. Scholarly requirements such as citations, captions, and cross-referencing can be enabled with a small amount of HTML or LaTeX and use of Pandoc.	http://daringfireball.net/projects/markdown/syntax
R markdown	A variant of markdown that extends it to allow chunks of R code to be embedded among the text. This results in a simple system for literate programming. For example, an R markdown document might have several paragraphs of text, then a chunk of R code that generates a figure, then several more paragraphs of text. Suitable for journal-article-length documents that include narrative text and output from statistical analysis.	http://rmarkdown.rstudio.com/
LaTeX	A complex document preparation system optimized for producing technical and scientific documentation. Suitable for large multi-part documents such as complex journal articles, books, and theses. Literate programming with R code interwoven among text is enabled <i>via</i> the knitr package.	https://latex-project.org
pandoc	An open source program for converting documents between a very wide variety of formats. Often used to convert markdown, R markdown and LaTeX documents to HTML (for web publication), PDF and Microsoft Word documents. It is built into RStudio.	http://pandoc.org/
Scientific programming		
script	A plain text file containing instructions for a computer written in a programming language, for example in R or Python.	
R	A free and open source programming language with strengths in data analysis and visualization. Most effective when used in combination with	https://www.r-project.org/

Table 1 (continued)

Term	Explanation	More information
	RStudio, a free and open source integrated development environment for R.	
Python	A free and open source programming language with a reputation for ease of use and being suitable for a wide range of scientific and commercial applications.	https://www.python.org/
MATLAB	A commercial programming language known for numerical and symbolic computing capabilities. The algorithms are proprietary, which means you cannot easily see the code of the algorithms and have to trust that MATLAB implemented it correctly. The proprietary nature also makes it hard, if not impossible, for others to extend or create tools for MATLAB.	http://www.mathworks.com/products/matlab
Version control		
Git	Open source software for version control and collaboration. It can handle any file type, but is most effective on plain text files such as scripts and markdown/LaTeX documents.	https://git-scm.com/
GitHub	A popular commercial web service that provides collaboration tools and free public hosting of files in git repositories. Private repositories are available for a fee. Similar services include GitLab and Bitbucket, both of which have the advantage of unlimited free private repositories.	https://github.com/
commit	A Git command to record changes in files to the Git repository. A sequence of commits creates a history of how the files have changed during your work on them.	http://git-scm.com/book/en/v2/Git-Basics-Recording-Changes-to-the-Repository
Computational environments		
Virtual machine	The use of software to emulate an entire operating system (such as GNU/Linux, Microsoft Windows, or Apple OS X) within another computer. For example, you might use a virtual machine to use a GNU/Linux operating system on a laptop where the main operating system is Microsoft Windows. Virtual machines are convenient for reproducing an entire computational environment, but they can consume a lot of hard disk space which makes sharing and archiving challenging.	
GNU/Linux	A free and open source computer operating system (<i>i.e.</i> , an alternative to Microsoft Windows and Apple OS X). Commonly used for scientific computing, internet servers, supercomputers, and Android phones and tablets. Popular distributions of GNU/Linux in academia include Ubuntu and Debian.	http://www.linux.org/
Linux container	A system for running multiple isolated Linux systems (containers) on a single Linux control host. Isolation means that the dependencies can	https://linuxcontainers.org/

Table 1 (continued)

Term	Explanation	More information
	be well understood and documented. In a research context, containers are useful for encapsulating the all of the diverse components of a complex data analysis system. Containers take up less disk space than a virtual machine, and so are more efficient for sharing and archiving.	
Docker	A free and open source system that simplifies the creation, use, sharing, and archiving of Linux containers. In a research context, Docker makes it easy to document and share computational environments so you can ensure that others have exactly the same software versions as you used.	https://www.docker.com/
Communities		
Software Carpentry	An international non-profit volunteer organization focusing on teaching researchers basic software skills. Prioritizes the use of free and open source software tools, encourages researchers to use permissive licenses for their research products. Target audience is novices with little or no prior computational experience.	http://software-carpentry.org/
Data Carpentry	Similar to Software Carpentry, but focuses more on domain-specific training covering the full lifecycle of data-driven research.	http://www.datacarpentry.org/
rOpenSci	A collaboration of volunteers from academia and industry developing R-based tools for making scientific research, data, and publication freely accessible to the public. They also conduct workshops to train researchers to use R and related tools.	https://ropensci.org/

further work is done without risk of being scooped. The data files should be archived in an online repository that issues persistent URLs (such as DOIs), that has a commitment to long-term sustainability (such as participation in the CLOCKSS scheme, Reich (2008)), and requires open licenses (such as CC-BY or CC-0) for datasets (Stodden 2009). Discipline-agnostic repositories include figshare.com and zenodo.org, and repositories and data sharing services specifically for archaeologists include the Archaeological Data Service, the Digital Archaeological Record, and Open Context (Arbuckle *et al.* 2014; Kansa *et al.* 2011).

Scripted Analyses

The dominant mode of interaction with data analysis tools for many researchers is a mouse-operated point-and-click interface with commercial software such as Microsoft's Excel, IBM's SPSS and SAS's JMP (Keeling and Pavur 2007; Thompson and Burnett 2012). This method of interaction is a formidable obstacle to reproducibility because mouse gestures leave few traces that are enduring and accessible to others (Wilson *et al.* 2014). *Ad hoc* edits of the raw data and analysis can easily

occur that leave no trace and interrupt the sequence of analytical steps (Sandve *et al.* 2013). While it is possible for a researcher to write down or even video their mouse-driven steps for others to reproduce, and this would be an excellent first step for sharing methods in many cases, these are rather cumbersome and inefficient methods for communicating many types of analyses. A second problem with much mouse-driven software is that the details of the data analysis are not available for inspection and modification because of the proprietary code of the software (Ince *et al.* 2012; Vihinen 2015). This constrains the transparency of research conducted with much commercial and mouse-driven software (Hatton and Roberts 1994).

While there are many conceivable methods to solve these problems (such as writing out all the operations in plain English or making a video screen-capture of the analysis), currently the most convenient and efficient solution is to interact with the data analysis tools using a script (Joppa *et al.* 2013). A script is a plain text file containing instructions composed in a programming language that direct a computer to accomplish a task. In a research context, researchers in fields such as physics, ecology, and biology write scripts to do data ingest, cleaning, analysis, visualizing, and reporting. By writing scripts, a very high-resolution record of the research workflow is created, and is preserved in a plain text file that can be reused and inspected by others (Gentleman and Temple Lang 2007). Data analysis using scripts has additional advantages of providing great flexibility to choose from a wide range of traditional and cutting-edge statistical algorithms, and tools for automation of repetitive tasks. Sharing these scripts may also increase the impact of the published research (Vandewalle 2012). The general approach of a scripted workflow to explicitly and unambiguously carry out instructions embodies the principles of reproducibility and transparency. Examples of programming languages used for scripting scientific analyses include R, Python, and MATLAB (Bassi 2007; Eglen 2009; Perkel 2015; Tippmann 2014). Among archaeologists who share code with their publications, R is currently the most widely used programming language (Bocinsky 2014; Bocinsky and Kohler 2014; Borck *et al.* 2015; Contreras and Meadows 2014; Crema *et al.* 2014; Drake *et al.* 2014; Dye 2011; Guedes *et al.* 2015; Lowe *et al.* 2014; Mackay *et al.* 2014; Marwick 2013; Peeples and Schachner 2012; Shennan *et al.* 2015).

Version Control

All researchers face the challenge of managing different versions of their computer files. A typical example, in the simple case of a solo researcher, is where multiple revisions of papers and datasets are saved as duplicate copies with slightly different file names (for example, appending the date to the end of the file name). In a more complex situation with multiple researchers preparing a report of publication, managing contributions from different authors and merging their work into a master document can result in a complex proliferation of files that can be very challenging to manage efficiently. While this complexity can be an inconvenience, it can lead to more profound problems of losing track of the provenance of certain results, and in the worst cases, losing track of the specific versions of files that produced the published results (Jones 2013).

One solution to these problems is to use a formal version control system (VCS) (Sandve *et al.* 2013), initially developed for managing contributions to large software

projects, and now used for many other purposes where multiple people are contributing to one file or collections of files. Instead of keeping multiple copies of a file, a VCS separately saves each change to a version control database (known as a “commit,” for example, the addition of a paragraph of text or a chunk of code) along with a comment describing the change. The commit history preserves a high-resolution record of the development of a file or set of files. Commits function as checkpoints where individual files or an entire project can be safely reverted to when necessary. Many VCSs allow for branching, where alternate ideas can be explored in a structured and documented way without disrupting the central flow of a project. Successful explorations can be merged into the main project, while dead ends can be preserved in an orderly way (Noble 2009). This is useful in two contexts, firstly to enable remote collaborators to work together without overwriting each other’s work and, secondly, to streamline responding questions from reviewers about why one option was chosen over another because all the analytical pathways explored by the authors are preserved in different branches in the VCS (Ram 2013). Version control is a key principle for reproducible research because of the transparency it provides. All decision points in the research workflow are explicitly documented so others can see why the project proceeded in the way it did. Researchers in many areas of science currently use Git or Subversion as a VCS (Jones 2013), often through a public or private online hosting service such as GitHub, BitBucket, or GitLab.

Computational Environments

Most researchers use one of three operating systems as their primary computational environment, Microsoft Windows, Apple OS X, or Linux. Once we look beyond the level of this basic detail, our computational environments diversify quickly, with many different versions of the same operating system in concurrent use, and many different versions of common data analysis software in concurrent use. For basic data analysis, the primary problem here is poor interoperability of file types from different versions of the same software. But for more complex projects that are dependent on several pieces of complex software from diverse sources, it is not uncommon for one of those pieces to change slightly (for example, when an update is released, a minor configuration is changed, or because different operating systems causes programs to behave differently), introducing unexpected output and possibly causing the entire workflow to fail (Glatard *et al.* 2015). For example, computationally intensive analyses often use mathematical functions based on single-precision floating-point arithmetic whose implementations vary between software (Keeling and Pavur 2007) and across operating systems. For archaeologists, this issue is particularly relevant to simulation studies. This situation can make it very challenging to create a research pipeline that will remain reproducible on any computer other than that of the researcher who constructed it (and into the future on the same computer, as its component software changes in ways that are beyond control of the researcher, due to automatic updates).

At the most general level, the principle that attempts to solve this problem is to provide a description of how other researchers can recreate the computational environment of the research pipeline. The simplest form of this is a list of the key pieces software and their version numbers; this is often seen in the archaeological literature where exotic algorithms are used. In other fields, where computationally intensive methods are more widespread,

and software dependencies are more extensive, more complex approaches have emerged, such as machine-readable instructions for recreating computational environments, or providing the entire actual computational environment that the analysis was conducted in (Dudley and Butte 2010; Howe 2012). Either of these provides another researcher with an identical copy of the operating systems and exact versions of all software dependencies. The ideal solution is to provide both, because providing the actual environment alone can result in a “black box” problem where the specific details of the environment are not available for inspection by another researcher, and the environment cannot easily be extended or joined to other environments for new projects. This results in a loss of transparency and portability, but this can be mitigated by providing a plain-text file that contains the instructions on how to recreate the environment in a machine-readable format. With this information, researchers can easily see the critical details of the environment, as well as efficiently recombine these details into other environments to create new research workflows. Examples of systems currently used by researchers to capture the entire environments include virtual machines (e.g., Oracle’s VirtualBox) and GNU/Linux containers (e.g., Docker). These environments are designed to be run in an existing operating system, so a researcher might have a GNU/Linux virtual machine running within their Windows or OS X computer. Vagrantfiles and Dockerfiles are common examples of machine-readable plain-text instructions for making virtual machines to an exact specification. One advantage of using self-contained computational environment like a virtual machine or container is that it is portable, and will perform identically whether it is used on the researcher’s laptop or high-performance facilities such as a commercial cloud computing service (Hoffa *et al.* 2008). While these more complex approaches may seem a bridge too far for most archaeologists, they offer some advantages for collaborating in a common computing environment (*i.e.*, in a project involving two or more computers using a virtual machine or container environment can simplify collaboration), and for working on small-scale iterations of an analysis prior to scaling up to time-consuming and expensive computations.

To summarize, in this section, I have described four general principles of reproducible research. These principles have been derived from current efforts to improve computational reproducibility in other fields, such as genomics, ecology, astronomy, climatology, neuroscience, and oceanography. The four principles are as follows: make data and code openly available and archive it in a suitable location, use a programming language to write scripts for data analysis and visualizations, use version control to manage multiple versions of files and contributions from collaborators, and, finally, document and share the computational environment of the analysis. Researchers following these principles will benefit from an increase in the transparency and efficiency of their research pipeline (Markowetz 2015). Results generated using these principles will be easier for other researchers to understand, reuse, and extend.

Case Study: The 1989 Excavation at Madjebebe, Northern Territory, Australia

In this section, I describe my efforts to produce a publication of archaeological research that demonstrates the above principles of reproducible research. I describe the specific tools that I used, explain my reasons for choosing these tools, and note any limitations and obstacles I encountered. Our paper on Madjebebe (Clarkson *et al.* 2015) describes

familiar types of evidence from a hunter-gatherer rockshelter excavation—stone artifacts, dates, sediments, mollusks. We—the co-authors of the Madjebebe paper and I—mostly used conventional and well-established methods of analyzing, summarizing, and visualizing the data. In this example, I expect the typical reader will recognize the types of raw data we used (measurements and observations from stone artifacts, dates, sediments, mollusks), and the output of our analysis (plots, tables, simple statistical test results). The novel component here is how we worked from the raw data to the published output. For this Madjebebe publication, we experimented with the principles of reproducible research outlined above, and used data archiving, a scripted analytical pipeline, version control, and an isolated computational environment. Additional details of our specific implementations are available at <https://github.com/benmarwick/1989-excavation-report-Madjebebe> and Marwick (2015).

That standard and familiar nature of the archaeological materials and methods used in the paper about Madjebebe should make it easy for the reader to understand how the methods for enhancing reproducibility described here can be adapted for the majority of research publications in archaeology. I recognize that not every research project can incorporate the use of these tools (for example, projects with very large amounts of data or very long compute times). However, my view is that the principles and tools described here are suitable for the majority of published research in archaeology (where datasets are small, *i.e.*, <10 GB, and analysis compute times are short, *i.e.*, <30 min).

Figshare for Data Archiving

We chose Figshare to archive all the files relating to the publication, including raw data, which we uploaded as a set of CSV files (Fig. 2). CSV stands for comma separated variables and is an open file format for spreadsheet files that can be opened and edited in any text editor or spreadsheet program. Although there are data repositories designed specifically for archaeologists (Beale 2012; Kansa 2012; Richards 1997), some of these are fee-based services and, at the time we deposited our data, they all lacked a programmatic interface and connections to other online services (such as GitHub, our version control backup service). Figshare is a commercial online digital repository

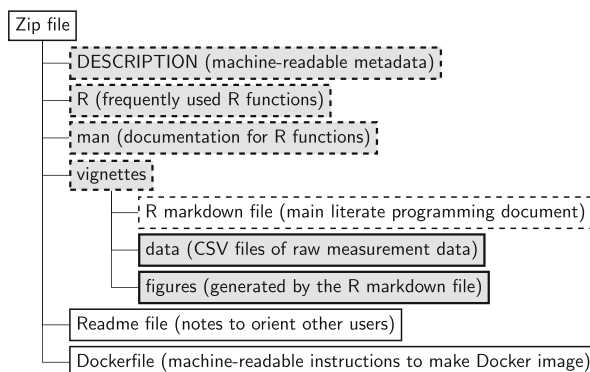


Fig. 2 File organization of the Figshare archive. The items with a dashed border are typical components of an R package, the solid outline indicates custom items added to form this specific compendium, and the shaded items indicate folders and the unshaded items indicate files

service that provides instant free unlimited archiving of any type of data files (up to 250 MB per file) for individual researchers in any field, and automatically issues persistent URLs (DOIs). Figshare also supplies file archiving services for many universities and publishers, including PLOS and Nature. Figshare allows the user to apply permissive Creative Commons licenses to archived files that specify how the files may be reused. We chose the CC0 license for our data files (equivalent to a release in the public domain); this is widely used and recommended for datasets (Stodden 2009). The CC0 license is simpler than the related CC-BY (requiring attribution) and CC-NC (prohibiting commercial use) license, so CC0 eliminates all uncertainty for potential users, encouraging maximal reuse and sharing of the data. We also archived our programming code on Figshare and applied the MIT license which is a widely used software license that permits any person to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the code (Henley and Kemp 2008; Morin *et al.* 2012). Our motivation for choosing these licenses is to clearly communicate to others that we are comfortable with our data and code to be reused in any way—with appropriate attributions resulting from normal scholarly practices (Stodden 2009). The MIT license has the added detail of specifically not providing a warranty of any kind and absolving us as authors from liability for any damages or problems that others might suffer or encounter when using our code.

R for Scripting the Analysis

I used the R programming language to script our data analysis and visualization workflow. I chose R because it is a highly expressive, functional, interpretive, object-oriented language that was originally developed by two academic statisticians in the 1990s (Chambers 2009; Wickham 2014). Like Python, R is a free and open source complete programming language. Where the two differ is that R is heavily customized for data analysis and visualization (Gandrud 2013b; Tippmann 2014). Python, which has a reputation for readability and ease of use, is a general-purpose programming tool with fewer customizations for data analysis and visualization (Perkel 2015). In the last decade, R has acquired a large user community of researchers, including archaeologists, many of whom contribute packages to a central open repository that extend the functionality of the language (Mair *et al.* 2015). These packages are typically accompanied by peer-reviewed scholarly publications that explain the algorithms presented in the package. Such a large and active community means that many common data analysis and visualization tasks have been greatly simplified by R packages, which is a key factor in my choice of this language. For example, rOpenSci is a collective of scientists mostly in ecology, evolution, and statistics that supports the development of R packages to access and analyze data, and provide training to researchers (Boettiger *et al.* 2015). Our publication depended on 19 of these user-contributed packages, which saved me a substantial amount of programming effort. I also organized our code as a custom R package because it provides a logical and widely shared structure to organizing the analysis and data files. The R package structure gives us access to the many quality control tools involved in package building, and is a convenient template for projects of any scale (Wickham 2015). Because packages are ubiquitous among R users, we hope that by providing our code as an R package, the use of familiar conventions for organizing the code will make it easier for other users to inspect, use, and extend our code.

The knitr and rmarkdown packages are especially relevant to our efforts to make our analysis reproducible (Xie 2013). Knitr provides algorithms for dynamically converting plain text and R code into formatted documents (*i.e.*, PDF, HTML, or MS Word) that contain the text and the output of the code, such as tables and plots. Rmarkdown provides an authoring format that enables the creation of dynamic documents using a simple syntax (related to HTML and LaTeX, but simpler) for formatting text and managing citations, captions, and other typical components of a scientific document (Baumer and Udwin 2015; Baumer *et al.* 2014). The rmarkdown package uses a document formatting language called markdown, which has a simple syntax for styling text, and extends it into a format called R markdown that enables embedded computation of R code contained in the markdown document. Using syntax for styling in markdown (and HTML, LaTeX, *etc.*) is different to composing and editing in Microsoft Word because markdown separates presentation from content. An example of this can be seen in the heading in Fig. 3, where the two hash symbols are the syntax for a heading, and the formatting is applied only when the document is executed. Together, the knitr and rmarkdown packages enabled us to compose a single plain-text source document that contained interwoven paragraphs of narrative text and chunks of R code. This approach has the code located in context with the text so any reader can easily see the role of the code in the narrative. This results in an executable paper (*cf.* Leisch *et al.* 2011; Nowakowski *et al.* 2011), which, when rendered by the computer using the knitr package, interprets the R code to generate the statistical and visual output and applies the formatting syntax to produce readable output in the form of a HTML, Microsoft Word, or PDF file that contains text, statistical results and tables, and data visualizations. This practice of having documentation and code in a single interwoven source document is known as literate programming (Knuth 1984). This is a focus of many efforts to improve the reproducibility of research, for example, by computer scientists

A Minimal Example of Literate Programming with R Markdown

We examined the relationship between artefact mass and distance from source using a linear regression model:

$$Y = \beta_0 + \beta_1 x + \epsilon$$

```


```{r echo=FALSE}
draw plot
plot(artefacts, pch = 20, col = 'darkgray')
#fit linear model
fit <- lm('mass (g)' ~ 'distance from source (km)',
 data = artefacts)
draw linear model regression line on plot
abline(fit, lwd = 2)
get slope value to use in text
slope <- round(coef(fit)[2],2)
```

```

The slope of a simple linear regression is `r slope`.

A Minimal Example of Literate Programming with R Markdown

We examined the relationship between artefact mass and distance from source using a linear regression model: $Y = \beta_0 + \beta_1 x + \epsilon$.



The slope of a simple linear regression is -4.1.

Fig. 3 A small literate programming example showing a sample of R markdown script similar to that used in our publication (on the *left*), and the rendered output (on the *right*). The example shows how to formulae can be included, and how a chunk of R code can be woven among narrative text. The code chunk draws a plot of artifact mass by distance from source, computes a linear regression, and adds the regression line to the plot. It also shows how one of the output values from the linear regression can be used in the narrative text without copying and pasting

and neuroscientists (Abari 2012; Delescluse *et al.* 2012; Schulte *et al.* 2012; Stanislac *et al.* 2015), but is not currently a mainstream practice in any field.

Git and GitHub for Version Control and Code Sharing

I chose Git as our version control system because it is currently by far the most widely used version control system at the moment, both in research contexts and for software engineering (Jones 2013; Loeliger and McCullough 2012). Git is a free and open source cross-platform program for tracking changes in plain text documents. The current popularity of Git is important because it means there is a lot of documentation and examples available to learn how to use the system. The key benefit of using Git was saving episodes of code-writing in meaningful units, for example the preparation of each figure was a single commit (Fig. 4). This was helpful because if some new code had an unexpected effect on an earlier figure, I could revert back to the previous commit where the code worked as expected. This high-resolution control over the progress of the code-writing provided by the version control system was helpful for identifying and solving problems in the analysis. During the peer-review and proofing stages, I used Git commits to indicate the exact version of the code that was used for the draft, revised, and final versions of the paper, which was helpful for keeping track of the changes we made in response to the reviewers' comments.

I used GitHub as a remote backup for our project, hosting the code and data files together with their Git database. GitHub is one of several commercial online services that hosts Git repositories and provides online collaboration tools (GitHub repositories that are open to the public are free, but fees are charged for private repositories; fee-waivers are available for academic users). While writing the paper, I worked on a private GitHub repository that was not publicly accessible because we needed approval from other stakeholders (such as the Aboriginal group on whose land the archaeological site is located) for the final paper before revealing it to the public. When the paper was published, I made the repository open and publicly available on GitHub (Barnes 2010), as well as archiving a copy of the code on Figshare with the data. The code on Figshare is frozen to match the output found in the published article, but the code on GitHub continues to be developed, mostly minor edits and improvements that do not change the content of the executed document. GitHub has Git-based tools for

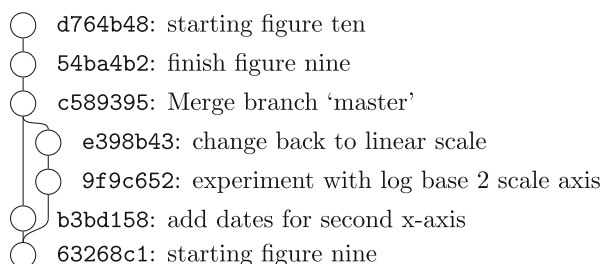


Fig. 4 Git commit history graph. This excerpt shows a typical sequence of commits and commit messages for a research project. The seven character code are keys that uniquely identify each commit. The example here shows the creation and merging of a branch to experiment with a variation of a plot axis. The graph shows more recent events at the top and earlier events at the bottom

organizing large-scale collaboration on research projects that are widely used in other fields, but I did not use these because of the small scale of our project (Gandrud 2013a).

Docker for Capturing the Computational Environment

Currently, there are two widely used methods for creating portable, isolated computational environments. The most established method is to create a virtual machine, usually taking the form of a common distribution of GNU/Linux such as Ubuntu or Debian. Although this is a widely used and understood method, it is also time-consuming to prepare the virtual machine, and the virtual machine occupies a relatively large amount of disk space (8 Gb in our case). I preferred the GNU/Linux container method because the virtual environment can be created much faster (which is more convenient for iteration) and the container image occupies much less disk space. The key difference between the two is that a virtual machine replicates an entire operating system, while the container image only shares some of the system resources to create an isolated computational environment, rather than requiring a complete system for each environment (Fig. 5). The low resource use of the container system makes it possible to




Fig. 5 Schematic of computer memory use of Docker (on the left) compared to a typical virtual machine (on the right). This figure shows how much more efficiently Docker uses hardware resources such as hard drive space, compared to a virtual machine

run several virtual environments simultaneously on a Windows or Mac desktop or laptop computer.

The specific GNU/Linux container system we used is called Docker, and is currently the dominant open source container system (Boettiger 2015). Like Git and R, Docker is a free and open source program. Docker is developed by a consortium of software companies, and they host an open, version-controlled online repository of ready-made Docker images, known as the Docker Hub, including several that contain R and RStudio in the GNU/Linux operating system. We used images provided by rOpenSci as our base image, and wrote a Dockerfile to specify further customizations on this base image. These include the installation of the JAGS library (Plummer and others 2003) to enable efficient Bayesian computation in R. Our Docker image is freely available on the Docker Hub and may be accessed by anyone wanting access to the original computational environment that we used for our analysis. Similarly, our Dockerfile is included in our code repository so that the exact contents of our Docker image are described (for example, in case the Docker Hub is unavailable, a researcher can rebuild our Docker image from the Dockerfile). Using the Dockerfile, our image can be reconstituted and extended for other purposes. We treated our Docker image as a disposable and isolated component, deleting and recreating it regularly to be sure that the computational environment documented in the Dockerfile could run our analyses.

Discussion

Developing competence in using these tools for enhancing computational reproducibility is time-consuming, and raises the question of how much of this is practical for most archaeologists, and what the benefits and costs might be. Our view is that once the initial costs of learning the tools is paid off, implementing the principles outlined above makes research and analysis easier, and has material professional benefits.

Perhaps the best established benefit is that papers with publicly available datasets receive a higher number of citations than similar studies without available data. Piwowar *et al.* (2007) investigated 85 publications on microarray data from clinical trials and found that papers that archived their data were cited 69 % more often than papers that did not archive. However, a larger follow-up study by Piwowar and Vision (2013) of 10,557 articles that created gene expression microarray data discovered only a 9 % citation advantage for papers with archived data. Henneken and Accomazzi (2011) analyzed 3814 articles in four astronomy journals and found that articles with links to open datasets on average acquired 20 % more citations than articles without links to data. Restricting the sample to papers published since 2009 in *The Astrophysical Journal*, Dorch (2012) found that papers with links to data are receiving 50 % more citations per paper per year, than papers without links to data. In 1331 articles published in *Paleoceanography* between 1993 and 2010, Sears (2011) found that publicly available data in articles was associated with a 35 % increase in citations. While I am not aware of any studies specifically of archaeological literature, similar positive effects of data sharing have been described in the social sciences. In 430 articles in the *Journal of Peace Research*, articles that offered data in any form, either through appendices, URLs, or contact addresses, were on average cited twice as frequently as an article with no data but otherwise equivalent author credentials and

article variables (Gleditsch and Strand 2003). It is clear that researchers in a number of different fields following the first principle of reproducible research benefit from a citation advantage for their articles that include publicly available datasets. In addition to increased citations for data sharing, Pienta *et al.* (2010) found that data sharing is associated with higher publication productivity. They examined 7040 NSF and NIH awards and concluded that a typical research grant award produces a median of five publications, but when data are archived a research grant award leads to a median of ten publications.

It is also worth noting that the benefits of using a programming language such as R for archaeological analyses extend beyond enhanced reproducibility. From a practical standpoint, users of R benefit from it being freely available for Windows, Unix systems (such as GNU/Linux), and the Mac. As a programming language designed for statistics and data visualization, R has the advantage of providing access to many more methods than commercial software packages such as Excel and SPSS. This is due to its status as the lingua franca for academic statisticians (Morandat *et al.* 2012; Narasimhan *et al.* 2005; Widemann *et al.* 2013), which means that R is the development environment for many recently developed algorithms found in journals (Bonhomme *et al.* 2014; Reshef *et al.* 2011), and these algorithms are readily available for archaeologists and others to use. R is widely known for its ability to complex data visualizations and maps with just a few lines of code (Bivand *et al.* 2008; Kahle and Wickham 2013; Sarkar 2008; Wickham 2009). Furthermore, my view is that once the learning curve is overcome, for most analyses using R would not take any longer than alternative technologies, and will often save time when previously written code is reused in new projects.

The primary cost of enhancing reproducibility is the time required to learn to use the software tools. I did not quantify this directly, but my personal experience is that about 3 years of self-teaching and daily use of R was necessary to develop the skills to code the entire workflow of our case study. Much less time was needed to learn Git and Docker, because the general concepts of interacting with these types of programs are similar to working with R (for example, using a command line interface and writing short functions using parameters). I expect that most archaeologists could develop competence substantially quicker than I did, by participating in short training courses such as those offered by Software Carpentry (Wilson 2014), Data Carpentry (Teal *et al.* 2015), rOpenSci (Boettiger *et al.* 2015), and similar organizations, or through the use of R in quantitative methods courses. I did not measure the amount of time required to improve the reproducibility of our case study article because I planned the paper to be reproducible before we started the analysis. This makes it difficult to separate time spent on analytical tasks from time spent on tasks specifically related to reproducibility. This situation, where the case study has “built-in reproducibility” and the additional time and effort is marginal, may be contrasted with “bolt-on reproducibility,” where reproducibility is enhanced only after the main analysis is complete. In the “bolt-on” situation, I might estimate a 50 % increase in the amount of time required for a project similar to this one. For multi-year projects with multiple teams, the time needed for the bolt-on approach would probably make it infeasible.

The main challenge I encountered using the tools described above in project was the uneven distribution of familiarity with them across our team. This meant that much of the final data analysis and visualization work presented in the publication was

concentrated on the team members familiar with these tools. The cause of this challenge is mostly likely the focus on point-and-click methods in most undergraduate courses on data analysis (Sharpe 2013). The absence of discussion of software in the key texts on statistics and archaeology (VanPool and Leonard 2010) is also a contributing factor. This contrasts with other fields where statistical methods and the computational tools to implement them are often described together (Buffalo 2015; Haddock and Dunn 2011; Scopatz and Huff 2015). This makes it difficult for archaeologists to acquire the computational skills necessary to enable reproducible research during a typical archaeology degree, leaving only self-teaching and short workshops as options for the motivated student.

Conclusion

We have outlined one potential standard way for enhancing the reproducibility of archaeological research, summarized in Fig. 1 and Table 2. Our compendium is a collection of files that follows the formal structure of an R package, and includes the raw data, R scripts organized into functions and an executable document, a Git database that includes the history of changes made to all the files in the compendium, and a Dockerfile that recreates the computational environment of our analysis. While the exact components of this kind of compendium will undoubtedly change over time as newer technologies appear, I expect that the general principles I have outlined will remain relevant long after these specific technologies have faded from use.

Two future directions follow from the principles, tools, and challenges that I have discussed above. First, the rarity of archaeologists with the computational skills necessary for reproducible research (as I observed on our group, and in the literature broadly, Table 2) highlights the need for future archaeologists to be trained as Pi-shaped researchers, rather than T-shaped researchers (Fig. 6). Current approaches to postgraduate training for archaeologists results in T-shaped researchers with wide-but-shallow general knowledge, but deep expertise and skill in one particular area. In contrast, a Pi-shaped researcher has the same wide breadth, but has to have deep knowledge of both their own domain-specific specialization, as well as a second area of deep knowledge in the computational principles and tools that enable reproducible research (Faris *et al.* 2011).

A second future direction is the need to incentivize training in, and practicing of, reproducible research by changing the editorial standards of archaeology journals. Although all the technologies and infrastructure to enhance research reproducibility are already available, they are not going to be widely used by researchers until there are strong incentives and a detailed mandate (McCullough and Vinod 2003; McCullough *et al.* 2006, 2008). One way to incentivize improvements to reproducibility is for journal editors to require submission of research compendia in place of the conventional stand-alone manuscript submission (Miguel *et al.* 2014). A research compendium is a manuscript accompanied by code and data files (or persistent links to reputable online repositories) that allows reviewers and readers to reproduce and extend the results without needing any further materials from the original authors (Gentleman and Temple Lang 2007; King 1995). This paper is an example of a research compendium, with the source files available at <http://dx.doi.org/10.6084/m9.figshare.1563661>, and

Table 2 Summary of degrees of reproducibility

| Degree of reproducibility | Data | Analysis | Computational environment | Comment |
|--------------------------------|--|---|---|---|
| Not reproducible | Summary statistics of the raw data are presented. | Brief narrative of methods is presented. | No information is provided. | The current status quo for scholarly journal articles. |
| Low reproducibility | The reader invited to contact the author for access to the data. | Brief narrative of methods is presented; names and version numbers of software are stated. | No information is provided. | Frequently seen. Inviting readers to contact the author to access the raw data is no guarantee that the raw data is available. |
| Low-Moderate reproducibility | The journal article is accompanied by files of raw data tables in PDF or Excel (<i>i.e.</i> , binary) files. | Brief narrative of methods is presented; names and version numbers of software are stated. | No information is provided. | Frequently seen. Having the raw data in supplementary material makes it much more accessible compared to when it must be requested from the author. However, extracting raw data from a PDF or other binary file format can be time-consuming and introduce errors. This presents obstacles to re-use of the data. |
| Moderate-High reproducibility | The journal article is accompanied by plain text files (<i>e.g.</i> , CSV format) of raw data. | The journal article is accompanied by script files of R or Python code that demonstrate key parts of the analysis (but do not generate all the results presented in the paper). | No information is provided. | Uncommon. Raw data in plain text format makes re-use highly efficient. Script files with code provide valuable insights into analytical decisions that are not narrated in the text of the article. However, because the code is not complete, substantial effort and skill is required by other researchers to reproduce the results of the article, and to re-use the code in new studies. This presents obstacles to re-use of the code. |
| High-Very high reproducibility | The journal article includes DOIs to an open access repository that contains plain text files (<i>e.g.</i> , CSV format) of raw data. | The open access repository linked to from the paper includes version-controlled R package or script files of R or Python code to reproduce all | The open access repository linked to from the paper includes a dockerfile that documents the computational environment of the published | Currently rarely seen. Other researchers should have a good chance to reproduce, re-use, and extend the published results with this combination of plain text data files, code that documents every analysis |

Table 2 (continued)

| Degree of reproducibility | Data | Analysis | Computational environment | Comment |
|---------------------------|------|---|--|--|
| | | of the analysis output and graphics in the article. | analysis, and a docker image that allows another person to use that environment. | and visualization in the paper, and details of the computational environment of the original analysis. Note that this does not guarantee complete and permanent reproducibility, but it gives the best odds we can currently provide. The use of an open access repository means that researchers can access the files even if they do not have a subscription to the journal, and ensures the availability of the files if the journal website changes. |




Fig. 6 T-shaped and Pi-shaped researchers

the case study paper on Madgebebe is more realistic and complex example of a compendium, online at <http://dx.doi.org/10.6084/m9.figshare.1297059>. Requiring submission of compendia instead of simply manuscripts is currently being experimented with by journals in other fields (e.g., *Quarterly Journal of Political Science*, *Biostatistics*) (Nosek *et al.* 2015; Peng 2009). The results of these experiments suggest that changing research communication methods and tools is a slow process, but they are valuable to find mistakes in submissions that are otherwise not obvious to reviewers, and they show that such changes to editorial expectations are possible without the journal being abandoned by researchers.

In archaeology, much progress has already been made in this direction by researchers using agent-based modeling. Archaeological publications that employ agent-based models often make available the complete code for their model in a repository such as OpenABM, which has successfully established community norms for documenting and disseminating computer code for agent-based models (Janssen *et al.* 2008). In archaeological publications, especially where a new method is presented, there is an urgent need to converge on similar community norms of sharing data and code in standardized formats. This will speed the adoption of new methods by reducing the effort needed to reverse-engineer the publication in order to adapt the new method to a new research problem. Most archaeologists will benefit from publications (their own and others') being reproducible, but attaining a high degree of reproducibility may not be possible for some publications. For example, only a low degree of reproducibility is possible for research that depends on sensitive data that cannot be made public, or research that depends on algorithms in specialized, expensive proprietary software (such as those provided by research instrument manufacturers). However, I believe that the majority of archaeological research publications have ample scope for substantial improvements in reproducibility. The technical problems are largely solved; the challenge now is to change the norms of the discipline to make high reproducibility a canonical attribute of high-quality scholarly work.

Software pervades every domain of research, and despite its importance in generating results, the choice of tools is very personal (Healy 2011), and archaeologists are given little guidance in the literature or during training. With this paper, I hope to begin a discussion on general principles and specific tools to improve the computational reproducibility of published archaeological research. This discussion is important because the choice of tools has ethical implications about the reliability of claims made

in publication. Tools that do not facilitate well-documented, transparent, portable, and reproducible data analysis workflows may, at best, result in irreproducible, unextendable research that does little to advance the discipline. At worst, they may conceal accidents or fraudulent behaviors that impede scientific advancement (Baggerly and Coombes 2009; Herndon *et al.* 2014; Laine *et al.* 2007; Lang 1993; Miller 2006).

Acknowledgments Thanks to Chris Clarkson, Mike Smith, Richard Fullagar, Lynley A. Wallis, Patrick Faulkner, Tiina Manne, Elspeth Hayes, Richard G. Roberts, Zenobia Jacobs, Xavier Carah, Kelsey M. Lowe, and Jacqueline Matthews for their cooperation with the JHE paper. Thanks to the Mirarr Senior Traditional Owners, and to our research partners, the Gundjeimhi Aboriginal Corporation, for granting permission to carry out the research that was published in the JHE paper, and led to this paper. Thanks to Kyle Bocinsky and Oliver Nakoinz for their helpful peer reviews and many constructive suggestions. This research was carried out as part of ARC Discovery Project DP110102864. This work was supported in part by the University of Washington eScience Institute, and especially benefited from the expertise of the Reproducibility and Open Science working group. An earlier version was presented at an International Neuroinformatics Coordinating Facility (INCF) meeting in December 2014 organised by Stephen Eglen, and benefited from discussion during that meeting. I am a contributor to the Software and Data Carpentry projects and the rOpenSci collective; beyond this, I declare that I have no conflict of interest.

References

- Abari, K. (2012). Reproducible research in speech sciences. *International Journal of Computer Science Issues*, 9(6), 43–52. Retrieved from <http://www.ijcsi.org/papers/IJCSI-9-6-2-43-52.pdf>
- Arbuckle, B. S., Kansa, S. W., Kansa, E., Orton, D., Çakırlar, C., Gourichon, L., & Würtenberger, D. (2014). Data sharing reveals complexity in the westward spread of domestic animals across neolithic Turkey. *PLoS One*, 9(6), e99845. doi:10.1371/journal.pone.0099845.
- Baggerly, K. A., & Coombes, K. R. (2009). Deriving chemosensitivity from cell lines: forensic bioinformatics and reproducible research in high-throughput biology. *The Annals of Applied Statistics*, 1309–1334.
- Barnes, N. (2010). Publish your computer code: it is good enough. *Nature News*, 467(7317), 753–753. doi:10.1038/467753a.
- Bassi, S. (2007). A primer on Python for life science researchers. *PLoS Computational Biology*, 3(11). doi:10.1371/journal.pcbi.0030199
- Baumer, B., & Udwin, D. (2015). R markdown. *Wiley Interdisciplinary Reviews: Computational Statistics*, 7(3), 167–177. doi:10.1002/wics.1348.
- Baumer, B., Cetinkaya-Rundel, M., Bray, A., Loi, L., & Horton, N. J. (2014). R markdown: Integrating a reproducible analysis tool into introductory statistics. *Technology Innovations in Statistics Education*, 8(1). Retrieved from <http://www.escholarship.org/uc/item/90b2f5xh>
- Beale, N. (2012). How community archaeology can make use of open data to achieve further its objectives. *World Archaeology*, 44(4), 612–633.
- Begley, C. G., & Ioannidis, J. P. A. (2015). Reproducibility in science improving the standard for basic and preclinical research. *Circulation Research*, 116(1), 116–126. doi:10.1161/CIRCRESAHA.114.303819.
- Bivand, R. S., Pebesma, E. J., Gomez-Rubio, V., & Pebesma, E. J. (2008). *Applied spatial data analysis with R* (Vol. 747248717). Springer.
- Bocinsky, R. K. (2014). Extrinsic site defensibility and landscape-based archaeological inference: an example from the northwest coast. *Journal of Anthropological Archaeology*, 35, 164–176.
- Bocinsky, R. K., & Kohler, T. A. (2014). A 2,000-year reconstruction of the rain-fed maize agricultural niche in the US southwest. *Nature Communications*, 5.
- Boettiger, C. (2015). An introduction to docker for reproducible research. *SIGOPS Operating System Review*, 49(1), 71–79. doi:10.1145/2723872.2723882.
- Boettiger, C., Hart, T., Chamberlain, S., & Ram, K. (2015). Building software, building community: lessons from the rOpenSci project. *Journal of Open Research Software*, 1(1), e8.

- Bonhomme, V., Picq, S., Gaucherel, C., & Claude, J. (2014). Momocs: outline analysis using r. *Journal of Statistical Software*, 56(13), 1–24. doi:10.18637/jss.v056.i13.
- Borck, L., Mills, B. J., Peeples, M. A., & Clark, J. J. (2015). Are social networks survival networks? An example from the late pre-Hispanic US southwest. *Journal of Archaeological Method and Theory*, 22(1), 33–57.
- Borgman, C. L. (2012). The conundrum of sharing research data. *Journal of the American Society for Information Science and Technology*, 63(6), 1059–1078.
- Buckheit, J. B., & Donoho, D. L. (1995). WaveLab and reproducible research. In A. Antoniadis & G. Oppenheim (Eds.), *Wavelets and statistics* (pp. 55–81). Springer New York. Retrieved from http://link.springer.com/chapter/10.1007/978-1-4612-2544-7_5
- Buffalo, V. (2015). *Bioinformatics data skills reproducible and robust research with open source tools*. CA: O'Reilly Media.
- Chambers, J. M. (2009). *Software for data analysis: Programming with R (1st ed. 2008. Corr. 2nd printing 2009 edition.)*. New York: Springer.
- Clarkson, C., Smith, M., Marwick, B., Fullagar, R., Wallis, L. A., Faulkner, P., & Florin, S. A. (2015). The archaeology, chronology and stratigraphy of Madjedbebe (Malakunanja II): a site in northern australia with early occupation. *Journal of Human Evolution*, 83, 46–64. doi:10.1016/j.jhevol.2015.03.014.
- Contreras, D. A., & Meadows, J. (2014). Summed radiocarbon calibrations as a population proxy: a critical evaluation using a realistic simulation approach. *Journal of Archaeological Science*, 52, 591–608.
- Crema, E., Edinborough, K., Kerig, T., & Shennan, S. (2014). An approximate bayesian computation approach for inferring patterns of cultural evolutionary change. *Journal of Archaeological Science*, 50, 160–170.
- Dafoe, A. (2014). Science deserves better: the imperative to share complete replication files. *PS: Political Science & Politics*, 47(01), 60–66. doi:10.1017/S104909651300173X.
- Delescluse, M., Franconville, R., Joucla, S., Lieury, T., & Pouzat, C. (2012). Making neurophysiological data analysis reproducible: why and how? *Journal of Physiology-Paris*, 106(3–4), 159–170. doi:10.1016/j.jphysparis.2011.09.011.
- Donoho, D. L., Maleki, A., Rahman, I. U., Shahram, M., & Stodden, V. (2009). Reproducible research in computational harmonic analysis. *Computing in Science & Engineering*, 11(1), 8–18. doi:10.1109/MCSE.2009.15.
- Dorch, S. (2012). On the citation advantage of linking to data: Astrophysics. Retrieved from <https://halshs.archives-ouvertes.fr/hprints-00714715/>
- Drake, W. H. A. H., Lee, B., & Wills, W. (2014). Strontium isotopes and the reconstruction of the Chaco regional system: evaluating uncertainty with bayesian mixing models. *PloS One*, 9(5), e95580. doi:10.1371/journal.pone.0095580.
- Dudley, J. T., & Butte, A. J. (2010). Reproducible in silico research in the era of cloud computing. *Nature Biotechnology*, 28(11), 1181–1185. doi:10.1038/nbt1110-1181.
- Dye, T. S. (2011). A model-based age estimate for polynesian colonization of Hawai'i. *Archaeology in Oceania*, 46(3), 130–138.
- Eglen, S. J. (2009). A quick guide to teaching R programming to computational biology students. *PLoS Computational Biology*, 5(8), e1000482. doi:10.1371/journal.pcbi.1000482.
- Faris, J., Kolker, E., Szalay, A., Bradlow, L., Deelman, E., Feng, W., & Kolker, E. (2011). Communication and data-intensive science in the beginning of the 21st century. *Omics: A Journal of Integrative Biology*, 15(4), 213–215.
- Gandrud, C. (2013a). Github: a tool for social data development and verification in the cloud. *The Political Methodologist*, 20(2), 7–16.
- Gandrud, C. (2013b). *Reproducible research with R and RStudio*. CRC Press.
- Gentleman, R., & Temple Lang, D. (2007). Statistical analyses and reproducible research. *Journal of Computational and Graphical Statistics*, 16(1), 1–23. doi:10.1198/106186007X178663.
- Glatard, T., Lewis, L. B., Ferreira da Silva, R., Adalat, R., Beck, N., Lepage, C., & Evans, A. C. (2015). Reproducibility of neuroimaging analyses across operating systems. *Frontiers in Neuroinformatics*, 9, 12. doi:10.3389/fninf.2015.00012.
- Gleditsch, N. P., & Strand, H. (2003). Posting your data: will you be scooped or will you be famous? *International Studies Perspectives*, 4(1), 72–107. doi:10.1111/1528-3577.04105.
- Guedes, J. d., Jin, G., & Bocinsky, R. K. (2015). The impact of climate on the spread of rice to north-eastern China: a new look at the data from shandong province. *PloS One*, 10(6), e0130430.
- Haddock, S. H. D., & Dunn, C. W. (2011). *Practical computing for biologists*. MA: Sinauer Associates Sunderland.

- Hatton, L., & Roberts, A. (1994). How accurate is scientific software? *IEEE Transactions on Software Engineering*, 20(10), 785–797. doi:10.1109/32.328993.
- Healy, K. (2011). Choosing your workflow applications. *The Political Methodologist*, 18(2), 9–18.
- Henley, M., & Kemp, R. (2008). Open source software: an introduction. *Computer Law & Security Review*, 24(1), 77–85.
- Henneken, E. A., & Accomazzi, A. (2011). Linking to data - effect on citation rates in astronomy. *CoRR, abs/1111.3618*. Retrieved from <http://arxiv.org/abs/1111.3618>
- Hemdon, T., Ash, M., & Pollin, R. (2014). Does high public debt consistently stifle economic growth? A critique of Reinhart and Rogoff. *Cambridge Journal of Economics*, 38(2), 257–279. doi:10.1093/cje/bet075.
- Hoffa, C., Mehta, G., Freeman, T., Deelman, E., Keahey, K., Berriman, B., & Good, J. (2008). On the use of cloud computing for scientific workflows. In *IEEE fourth international conference on eScience, 2008. eScience'08* (pp. 640–645). doi:10.1109/eScience.2008.167
- Howe, B. (2012). Virtual appliances, cloud computing, and reproducible research. *Computing in Science & Engineering*, 14(4), 36–41. doi:10.1109/MCSE.2012.62.
- Ince, D. C., Hatton, L., & Graham-Cumming, J. (2012). The case for open computer programs. *Nature*, 482(7386), 485–488. doi:10.1038/nature10836.
- Janssen, M. A., Alessa, L. N., Barton, M., Bergin, S., & Lee, A. (2008). Towards a community framework for agent-based modelling. *Journal of Artificial Societies and Social Simulation*, 11(2), 6.
- Jones, Z. M. (2013). Git/GitHub, transparency, and legitimacy in quantitative research. *The Political Methodologist*, 21(1), 6–7. Retrieved from <http://zmjones.com/static/papers/git.pdf>
- Joppa, L. N., McNerny, G., Harper, R., Salido, L., Takeda, K., O'Hara, K., & Emmott, S. (2013). Troubling trends in scientific software use. *Science*, 340(6134), 814–815.
- Kahle, D., & Wickham, H. (2013). Ggmap: spatial visualization with ggplot2. *The R Journal*, 5(1), 144–161.
- Kansa, E. (2012). Openness and archaeology's information ecosystem. *World Archaeology*, 44(4), 498–520.
- Kansa, E. C., Kansa, S. W., & Watrall, E. (2011). Archaeology 2.0: new approaches to communication and collaboration. *Cotsen Digital Archaeology Series*.
- Keeling, K. B., & Pavur, R. J. (2007). A comparative study of the reliability of nine statistical software packages. *Computational Statistics & Data Analysis*, 51(8), 3811–3831. doi:10.1016/j.csda.2006.02.013.
- King, G. (1995). Replication, replication. *PS: Political Science & Politics*, 28(03), 444–452.
- Kintigh, K. (2006). The promise and challenge of archaeological data integration. *American Antiquity*, 71(3), 567–578. doi:10.2307/40035365.
- Kintigh, K. W., Altschul, J. H., Beaudry, M. C., Drennan, R. D., Kinzig, A. P., & Kohler, T. A. (2014). Grand challenges for archaeology. *Proceedings of the National Academy of Sciences*, 111(3), 879–880.
- Knuth, D. E. (1984). Literate programming. *The Computer Journal*, 27(2), 97–111.
- Laine, C., Goodman, S. N., Griswold, M. E., & Sox, H. C. (2007). Reproducible research: moving toward research the public can really trust. *Annals of Internal Medicine*, 146(6), 450–453. doi:10.7326/0003-4819-146-6-200703200-00154.
- Lang, S. (1993). Questions of scientific responsibility: the Baltimore case. *Ethics & Behavior*, 3(1), 3–72. doi:10.1207/s15327019eb0301_1.
- Leisch, F., Eugster, M., & Hothorn, T. (2011). Executable papers for the R community: the R2 platform for reproducible research. *Procedia Computer Science*, 4, 618–626.
- Loeliger, J., & McCullough, M. (2012). *Version control with Git: Powerful tools and techniques for collaborative software development*. "O'Reilly Media, Inc."
- Lowe, K. M., Wallis, L. A., Pardoe, C., Marwick, B., Clarkson, C., Manne, T., & Fullagar, R. (2014). Ground-penetrating radar and burial practices in western Arnhem Land, Australia. *Archaeology in Oceania*, 49(3), 148–157.
- Mackay, A., Sumner, A., Jacobs, Z., Marwick, B., Bluff, K., & Shaw, M. (2014). Putslaagte 1 (PL1), the Doring river, and the later middle stone age in southern Africa's winter rainfall zone. *Quaternary International*, 350, 43–58.
- Mair, P., Hofmann, E., Gruber, K., Hatzinger, R., Zeileis, A., & Hornik, K. (2015). Motivation, values, and work design as drivers of participation in the r open source project for statistical computing. *Proceedings of the National Academy of Sciences*, 112(48), 14788–14792. doi:10.1073/pnas.1506047112.
- Markowitz, F. (2015). Five selfish reasons to work reproducibly. *Genome Biology*, 16.
- Marwick, B. (2013). Multiple optima in Hoabinhian flaked stone artefact palaeoeconomics and palaeoecology at two archaeological sites in northwest Thailand. *Journal of Anthropological Archaeology*, 32(4), 553–564.
- Marwick, B. (2015). *Code and data repository for a report on the 1989 excavations at Madjebebe, northern territory, Australia*. Retrieved May 30, 2015, from <http://dx.doi.org/10.6084/m9.figshare.1297059>

- McCullough, B. (2007). Got replicability? The _journal of money, credit and banking_ archive. *Econ Journal Watch*, 4(3), 326–337.
- McCullough, B., & Vinod, H. D. (2003). Verifying the solution from a nonlinear solver: a case study. *American Economic Review*, 93(3), 873–892. doi:10.1257/000282803322157133.
- McCullough, B., McGeary, K. A., & Harrison, T. D. (2006). Lessons from the JMCB archive. *Journal of Money, Credit, and Banking*, 38(4), 1093–1107. doi:10.1353/mcb.2006.0061.
- McCullough, B., McGeary, K. A., & Harrison, T. D. (2008). Do economics journal archives promote replicable research? *The Canadian Journal of Economics / Revue Canadienne d'Economie*, 41(4), 1406–1420.
- Miguel, E., Camerer, C., Casey, K., Cohen, J., Esterling, K. M., Gerber, A., & Van der Laan, M. (2014). Promoting transparency in social science research. *Science (New York, N.Y.)*, 343(6166), 30–31. doi:10.1126/science.1245317.
- Miller, G. (2006). A scientist's nightmare: software problem leads to five retractions. *Science*, 314(5807), 1856–1857. doi:10.1126/science.314.5807.1856.
- Morandat, F., Hill, B., Osvald, L., & Vitek, J. (2012). Evaluating the design of the R language. In *ECOOP 2012—Object-oriented programming* (pp. 104–131). Springer.
- Morin, A., Urban, J., & Sliz, P. (2012). A quick guide to software licensing for the scientist-programmer. *PLoS Computational Biology*, 8(7), e1002598.
- Narasimhan, B., et al. (2005). Lisp-stat to Java to R. *Journal of Statistical Software*, 13(4), 1–10.
- Noble, W. S. (2009). A quick guide to organizing computational biology projects. *PLoS Computational Biology*, 5(7), e1000424. doi:10.1371/journal.pcbi.1000424.
- Nosek, B., Alter, G., Banks, G., Borsboom, D., Bowman, S., & Breckler, S. (2015). Promoting an open research culture: author guidelines for journals could help to promote transparency, openness, and reproducibility. *Science (New York, NY)*, 348(6242), 1422.
- Nowakowski, P., Ciepiela, E., Harzłak, D., Kocot, J., Kasztelnik, M., Bartyński, T., & Malawski, M. (2011). The collage authoring environment. *Procedia Computer Science*, 4, 608–617.
- Peebles, M. A., & Schachner, G. (2012). Refining correspondence analysis-based ceramic seriation of regional data sets. *Journal of Archaeological Science*, 39(8), 2818–2827.
- Peng, R. D. (2009). Reproducible research and biostatistics. *Biostatistics*, 10(3), 405–408. doi:10.1093/biostatistics/kxp014.
- Peng, R. D. (2011). Reproducible research in computational science. *Science (New York, Ny)*, 334(6060), 1226.
- Perkel, J. M. (2015). Programming: pick up Python. *Nature*, 518(7537), 125–126. doi:10.1038/518125a.
- Pienta, A. M., Alter, G. C., & Lyle, J. A. (2010). The enduring value of social science research: the use and reuse of primary research data.
- Piowar, H. A., & Vision, T. J. (2013). Data reuse and the open data citation advantage. *Peer Journal*, 1, e175.
- Piowar, H. A., Day, R. S., & Fridsma, D. B. (2007). Sharing detailed research data is associated with increased citation rate. *PloS One*, 2(3), e308. doi:10.1371/journal.pone.0000308.
- Plummer, M., & others. (2003). JAGS: A program for analysis of Bayesian graphical models using gibbs sampling. In *Proceedings of the 3rd international workshop on distributed statistical computing* (Vol. 124, p. 125). Technische Universit at Wien.
- Ram, K. (2013). Git can facilitate greater reproducibility and increased transparency in science. *Source Code for Biology and Medicine*, 8(1), 7.
- Reich, V. (2008). CLOCKSS—it takes a community. *The Serials Librarian*, 54(1-2), 135–139.
- Reshef, D. N., Reshef, Y. A., Finucane, H. K., Grossman, S. R., McVean, G., Turnbaugh, P. J., & Sabeti, P. C. (2011). Detecting novel associations in large data sets. *Science*, 334(6062), 1518–1524. doi:10.1126/science.1205438.
- Richards, J. D. (1997). Preservation and re-use of digital data: the role of the archaeology data service. *Antiquity*, 71(274), 1057–1057.
- Rieth, C. (2013). Report from the SAA board of directors. *SAA Archaeological Record*, May, 42–44.
- Sandve, G. K., Nekrutenko, A., Taylor, J., & Hovig, E. (2013). Ten simple rules for reproducible computational research. *PLoS Computational Biology*, 9(10), e1003285. doi:10.1371/journal.pcbi.1003285.
- Sarkar, D. (2008). *Lattice: Multivariate data visualization with R*. Springer Science & Business Media.
- Schulte, E., Davison, D., Dye, T., & Dominik, C. (2012). A multi-language computing environment for literate programming and reproducible research. *Journal of Statistical Software*, 46(3), 1–24. Retrieved from <http://yuyang0.github.io/static/doc/babel.pdf>
- Schwab, M., Karrenbach, M., & Claerbout, J. (2000). Making scientific computations reproducible. *Computing in Science & Engineering*, 2(6), 61–67.

- Scopatz, A., & Huff, K. D. (2015). *Effective computation in physics: Field guide to research with python*. CA: O'Reilly Media.
- Sears, J. (2011). Data sharing effect on article citation rate in paleoceanography. In *AGU fall meeting abstracts* (Vol. 1, p. 1628).
- Sharpe, D. (2013). Why the resistance to statistical innovations? Bridging the communication gap. *Psychological Methods*, 18(4), 572. Retrieved from <http://psycnet.apa.org/journals/met/18/4/572/>
- Shennan, S. J., Crema, E. R., & Kerig, T. (2015). Isolation-by-distance, homophily, and 'core' vs 'package' cultural evolution models in neolithic europe. *Evolution and Human Behavior*, 36(2), 103–109.
- Stanisic, L., Legrand, A., & Danjean, V. (2015). An effective git and org-mode based workflow for reproducible research. *ACM SIGOPS Operating Systems Review*, 49(1), 61–70. Retrieved from <http://dl.acm.org/citation.cfm?id=2723881>
- Stodden, V. (2009). The legal framework for reproducible scientific research: licensing and copyright. *Computing in Science & Engineering*, 11(1), 35–40. doi:10.1109/MCSE.2009.19.
- Stodden, V., & Miguez, S. (2014). Best practices for computational science: software infrastructure and environments for reproducible and extensible research. *Journal of Open Research Software*, 2(1), e21. doi:10.5334/jors.ay.
- Stodden, V., Guo, P., & Ma, Z. (2013). Toward reproducible computational research: an empirical analysis of data and code policy adoption by journals. *PloS One*, 8(6), e67111. doi:10.1371/journal.pone.0067111.
- Teal, T. K., Cranston, K. A., Lapp, H., White, E., Wilson, G., Ram, K., & Pawlik, A. (2015). Data carpentry: workshops to increase data literacy for researchers. *International Journal of Digital Curation*, 10(1), 135–143.
- Tenopir, C., Allard, S., Douglass, K., Aydinoglu, A. U., Wu, L., Read, E., & Frame, M. (2011). Data sharing by scientists: practices and perceptions. *PloS One*, 6(6), e21101. doi:10.1371/journal.pone.0021101.
- Thompson, P. A., & Burnett, A. (2012). Reproducible research. *CORE Issues in Professional and Research Ethics*, 1(6).
- Tippmann, S. (2014). Programming tools: adventures with R. *Nature*, 517(7532), 109–110. doi:10.1038/517109a.
- Vandewalle, P. (2012). Code sharing is associated with research impact in image processing. *Computing in Science and Engineering*, 14(4), 42–47.
- VanPool, T. L., & Leonard, R. D. (2010). *Quantitative analysis in archaeology (1 edition.)*. Chichester, West Sussex, U.K. Malden, MA: Wiley-Blackwell.
- Vihinen, M. (2015). No more hidden solutions in bioinformatics. *Nature*, 521(7552), 261–261. doi:10.1038/521261a.
- Wicherts, J. M., Bakker, M., & Molenaar, D. (2011). Willingness to share research data is related to the strength of the evidence and the quality of reporting of statistical results. *PloS One*, 6(11), e26828. doi:10.1371/journal.pone.0026828.
- Wickham, H. (2009). *Ggplot2: Elegant graphics for data analysis*. Springer Science & Business Media.
- Wickham, H. (2014). *Advanced R (1 edition.)*. Boca Raton, FL: Chapman; Hall/CRC.
- Wickham, H. (2015). *R packages (1 edition.)*. Sebastopol, Calif.: O'Reilly Media.
- Widemann, B. T. y, Bolz, C. F., & Grelck, C. (2013). The functional programming language r and the paradigm of dynamic scientific programming. In H.-W. Loidl & R. Peña (Eds.), *Trends in functional programming* (Vol. 7829, pp. 182–197). Springer Berlin Heidelberg. doi:10.1007/978-3-642-40447-4_12
- Wilson, G. (2014). Software carpentry: Lessons learned. *F1000Research*, Retrieved from <http://f1000research.com/articles/3-62/v1>.
- Wilson, G., Aruliah, D. A., Brown, C. T., Chue Hong, N. P., Davis, M., Guy, R. T., & Wilson, P. (2014). Best practices for scientific computing. *PloS Biology*, 12(1), e1001745. doi:10.1371/journal.pbio.1001745.
- Xie, Y. (2013). *Dynamic documents with R and knitr*. CRC Press.